

softmax.py

```

1 import numpy as np
2
3 class Softmax(object):
4
5     def __init__(self, dims=[10, 3073]):
6         self.init_weights(dims=dims)
7
8     def init_weights(self, dims):
9         """
10        Initializes the weight matrix of the Softmax classifier.
11        Note that it has shape (C, D) where C is the number of
12        classes and D is the feature size.
13        """
14        self.W = np.random.normal(size=dims) * 0.0001
15
16    def loss(self, X, y):
17        """
18        Calculates the softmax loss.
19
20        Inputs have dimension D, there are C classes, and we operate on minibatches
21        of N examples.
22
23        Returns a tuple of:
24        - loss as single float
25        """
26
27
28        # Initialize the loss to zero.
29        loss = 0.0
30
31
32        # ===== #
33        # YOUR CODE HERE:
34        # Calculate the normalized softmax loss. Store it as the variable loss.
35        # (That is, calculate the sum of the losses of all the training
36        # set margins, and then normalize the loss by the number of
37        # training examples.)
38        # ===== #
39        minibatch_size = y.shape[0]
40        input_scores = self.W.dot(X.T)
41        loss = (1/float(minibatch_size)) * np.sum(np.log(np.sum(np.exp(input_scores), axis=0))) - np.choose(y, input_scores)
42        # ===== #
43
44        # END YOUR CODE HERE
45        # ===== #
46
47
48    return loss
49
50
51    def loss_and_grad(self, X, y):
52        """
53        Same as self.loss(X, y), except that it also returns the gradient.
54
55        Output: grad -- a matrix of the same dimensions as W containing
56        the gradient of the loss with respect to W.
57        """
58
59        # Initialize the loss and gradient to zero.
60        loss = 0.0
61        grad = np.zeros_like(self.W)
62
63        # ===== #
64        # YOUR CODE HERE:
65        # Calculate the softmax loss and the gradient. Store the gradient
66        # as the variable grad.
67        # ===== #
68        minibatch_size = y.shape[0]
69        input_scores = self.W.dot(X.T)
70        softmax_nominators = np.exp(self.W.dot(X.T))
71        softmax_denominators = np.sum(softmax_nominators, axis=0)
72        loss = (1 / float(minibatch_size)) * np.sum(
73            np.log(np.sum(np.exp(input_scores), axis=0))) - np.choose(y, input_scores)
74        for i in range(self.W.shape[0]):
75            for k in range(X.shape[0]):
76                if y[k] == i:
77                    grad[i] += X[k] * (softmax_nominators[i][k]/softmax_denominators[k] - 1)
78                else:
79                    grad[i] += X[k] * (softmax_nominators[i][k] / softmax_denominators[k])
80            grad[i] /= minibatch_size
81        # ===== #
82        # END YOUR CODE HERE
83        # ===== #
84
85    return loss, grad
86
87    def grad_check_sparse(self, X, y, your_grad, num_checks=10, h=1e-5):
88        """
89        sample a few random elements and only return numerical
90        in these dimensions.
91        """
92
93        for i in np.arange(num_checks):
94            ix = tuple([np.random.randint(m) for m in self.W.shape])

```

```

95     oldval = self.W[ix]
96     self.W[ix] = oldval + h # increment by h
97     fxph = self.loss(X, y)
98     self.W[ix] = oldval - h # decrement by h
99     fmxh = self.loss(X,y) # evaluate f(x - h)
100    self.W[ix] = oldval # reset
101
102    grad_numerical = (fxph - fmxh) / (2 * h)
103    grad_analytic = your_grad[ix]
104    rel_error = abs(grad_numerical - grad_analytic) / (abs(grad_numerical) + abs(grad_analytic))
105    print('numerical: %f analytic: %f, relative error: %e' % (grad_numerical, grad_analytic, rel_error))
106
107 def fast_loss_and_grad(self, X, y):
108 """
109     A vectorized implementation of loss_and_grad. It shares the same
110     inputs and ouputs as loss_and_grad.
111 """
112
113 loss = 0.0
114 grad = np.zeros(self.W.shape) # initialize the gradient as zero
115
116 # ===== #
117 # YOUR CODE HERE:
118 #   Calculate the softmax loss and gradient WITHOUT any for loops.
119 # ===== #
120 minibatch_size = y.shape[0]
121 num_features = X.shape[1]
122 input_scores = self.W.dot(X.T)
123 loss = (1 / float(minibatch_size)) * np.sum(
124     np.log(np.sum(np.exp(input_scores - np.amax(input_scores, axis=0)), axis=0)) - np.choose(y, input_scores - np.amax(input_scores, axis=0)))
125 softmax_nominators = np.exp(input_scores - np.amax(input_scores, axis=0))
126 softmax_denominators = np.sum(softmax_nominators, axis=0)
127 softmax_matrix = softmax_nominators / softmax_denominators
128 softmax_matrix[y, np.arange(minibatch_size)] -= 1
129 softmax_matrix = np.tile(softmax_matrix.T, (1, 1, 1))
130 grad = (1/float(minibatch_size)) * np.sum(softmax_matrix.T * X, axis=1)
131 # ===== #
132 # END YOUR CODE HERE
133 # ===== #
134
135 return loss, grad
136
137 def train(self, X, y, learning_rate=1e-3, num_iters=100,
138           batch_size=200, verbose=False):
139 """
140     Train this linear classifier using stochastic gradient descent.
141
142     Inputs:
143     - X: A numpy array of shape (N, D) containing training data; there are N
144       training samples each of dimension D.
145     - y: A numpy array of shape (N,) containing training labels; y[i] = c
146       means that X[i] has label 0 <= c < C for C classes.
147     - learning_rate: (float) learning rate for optimization.
148     - num_iters: (integer) number of steps to take when optimizing
149     - batch_size: (integer) number of training examples to use at each step.
150     - verbose: (boolean) If true, print progress during optimization.
151
152     Outputs:
153     A list containing the value of the loss function at each training iteration.
154 """
155     num_train, dim = X.shape
156     num_classes = np.max(y) + 1 # assume y takes values 0...K-1 where K is number of classes
157
158     self.init_weights(dims=[np.max(y) + 1, X.shape[1]]) # initializes the weights of self.W
159
160     # Run stochastic gradient descent to optimize W
161     loss_history = []
162
163     for it in np.arange(num_iters):
164         X_batch = None
165         y_batch = None
166
167         # ===== #
168         # YOUR CODE HERE:
169         #   Sample batch_size elements from the training data for use in
170         #   gradient descent. After sampling,
171         #   - X_batch should have shape: (dim, batch_size)
172         #   - y_batch should have shape: (batch_size,)
173         #   The indices should be randomly generated to reduce correlations
174         #   in the dataset. Use np.random.choice. It's okay to sample with
175         #   replacement.
176         # ===== #
177         idx = np.random.randint(low=0, high=X.shape[0], size=batch_size)
178         X_batch = X[idx]
179         y_batch = y[idx]
180
181         # ===== #
182         # END YOUR CODE HERE
183         # ===== #
184
185         # evaluate loss and gradient
186         loss, grad = self.fast_loss_and_grad(X_batch, y_batch)
187         loss_history.append(loss)
188
189         # ===== #
190         # YOUR CODE HERE:
191         #   Update the parameters, self.W, with a gradient step
192         # ===== #
193         self.W = self.W - learning_rate * grad

```

```
193
194     # ===== #
195     # END YOUR CODE HERE
196     # ===== #
197
198     if verbose and it % 100 == 0:
199         print('iteration {} / {}: loss {}'.format(it, num_iters, loss))
200
201     return loss_history
202
203 def predict(self, X):
204     """
205     Inputs:
206     - X: N x D array of training data. Each row is a D-dimensional point.
207
208     Returns:
209     - y_pred: Predicted labels for the data in X. y_pred is a 1-dimensional
210       array of length N, and each element is an integer giving the predicted
211       class.
212     """
213     y_pred = np.zeros(X.shape[1])
214     # ===== #
215     # YOUR CODE HERE:
216     #   Predict the labels given the training data.
217     # ===== #
218     y_pred = np.argmax(self.W.dot(X.T), axis=0)
219     # ===== #
220     # END YOUR CODE HERE
221     # ===== #
222
223     return y_pred
224
225
```