

## fc\_net.py

```

1  import numpy as np
2
3  from .layers import *
4  from .layer_utils import *
5
6  """
7  This code was originally written for CS 231n at Stanford University
8  (cs231n.stanford.edu). It has been modified in various areas for use in the
9  ECE 239AS class at UCLA. This includes the descriptions of what code to
10 implement as well as some slight potential changes in variable names to be
11 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung for
12 permission to use this code. To see the original version, please visit
13 cs231n.stanford.edu.
14 """
15
16 class TwoLayerNet(object):
17 """
18 A two-layer fully-connected neural network with ReLU nonlinearity and
19 softmax loss that uses a modular layer design. We assume an input dimension
20 of D, a hidden dimension of H, and perform classification over C classes.
21
22 The architecture should be affine - relu - affine - softmax.
23
24 Note that this class does not implement gradient descent; instead, it
25 will interact with a separate Solver object that is responsible for running
26 optimization.
27
28 The learnable parameters of the model are stored in the dictionary
29 self.params that maps parameter names to numpy arrays.
30 """
31
32 def __init__(self, input_dim=3*32*32, hidden_dims=100, num_classes=10,
33             dropout=0, weight_scale=1e-3, reg=0.0):
34 """
35 Initialize a new network.
36
37 Inputs:
38 - input_dim: An integer giving the size of the input
39 - hidden_dims: An integer giving the size of the hidden layer
40 - num_classes: An integer giving the number of classes to classify
41 - dropout: Scalar between 0 and 1 giving dropout strength.
42 - weight_scale: Scalar giving the standard deviation for random
43   initialization of the weights.
44 - reg: Scalar giving L2 regularization strength.
45 """
46 self.params = {}
47 self.reg = reg
48
49 # ===== #
50 # YOUR CODE HERE:
51 #   Initialize W1, W2, b1, and b2. Store these as self.params['W1'],
52 #   self.params['W2'], self.params['b1'] and self.params['b2']. The
53 #   biases are initialized to zero and the weights are initialized
54 #   so that each parameter has mean 0 and standard deviation weight_scale.
55 #   The dimensions of W1 should be (input_dim, hidden_dim) and the
56 #   dimensions of W2 should be (hidden_dims, num_classes)
57 # ===== #
58
59 self.params['W1'] = weight_scale * np.random.randn(input_dim, hidden_dims)
60 self.params['b1'] = np.zeros(hidden_dims)
61 self.params['W2'] = weight_scale * np.random.randn(hidden_dims, num_classes)
62 self.params['b2'] = np.zeros(num_classes)
63
64 # ===== #
65 # END YOUR CODE HERE
66 # ===== #
67
68 def loss(self, X, y=None):
69 """
70 Compute loss and gradient for a minibatch of data.
71
72 Inputs:
73 - X: Array of input data of shape (N, d_1, ..., d_k)
74 - y: Array of labels, of shape (N,). y[i] gives the label for X[i].
75
76 Returns:
77 If y is None, then run a test-time forward pass of the model and return:
78 - scores: Array of shape (N, C) giving classification scores, where
79   scores[i, c] is the classification score for X[i] and class c.

```

```

80
81     If y is not None, then run a training-time forward and backward pass and
82     return a tuple of:
83     - loss: Scalar value giving the loss
84     - grads: Dictionary with the same keys as self.params, mapping parameter
85             names to gradients of the loss with respect to those parameters.
86     """
87     scores = None
88
89     # ===== #
90     # YOUR CODE HERE:
91     #   Implement the forward pass of the two-layer neural network. Store
92     #   the class scores as the variable 'scores'. Be sure to use the layers
93     #   you prior implemented.
94     # ===== #
95
96     layer_1_out, layer_1_cache = affine_relu_forward(np.array(X), self.params['W1'], self.params['b1'])
97     scores, layer_2_cache = affine_forward(np.array(layer_1_out), self.params['W2'], self.params['b2'])
98
99     # ===== #
100    # END YOUR CODE HERE
101    # ===== #
102
103    # If y is None then we are in test mode so just return scores
104    if y is None:
105        return scores
106
107    loss, grads = 0, {}
108    # ===== #
109    # YOUR CODE HERE:
110    #   Implement the backward pass of the two-layer neural net. Store
111    #   the loss as the variable 'loss' and store the gradients in the
112    #   'grads' dictionary. For the grads dictionary, grads['W1'] holds
113    #   the gradient for W1, grads['b1'] holds the gradient for b1, etc.
114    #   i.e., grads[k] holds the gradient for self.params[k].
115    #
116    #   Add L2 regularization, where there is an added cost 0.5*self.reg*W^2
117    #   for each W. Be sure to include the 0.5 multiplying factor to
118    #   match our implementation.
119    #
120    #   And be sure to use the layers you prior implemented.
121    # ===== #
122
123    W1_norm_sqrd = (np.linalg.norm(self.params['W1']) ** 2)
124    W2_norm_sqrd = (np.linalg.norm(self.params['W2']) ** 2)
125
126    loss, soft_grad = softmax_loss(scores, y)
127    loss += 0.5 * self.reg * (W1_norm_sqrd + W2_norm_sqrd)
128    layer_2_back_grad, grads['W2'], grads['b2'] = affine_backward(soft_grad, layer_2_cache)
129    layer_1_back_grad, grads['W1'], grads['b1'] = affine_relu_backward(layer_2_back_grad, layer_1_cache)
130
131    grads['W2'] += self.reg * self.params['W2']
132    grads['W1'] += self.reg * self.params['W1']
133    # ===== #
134    # END YOUR CODE HERE
135    # ===== #
136
137    return loss, grads
138
139
140    class FullyConnectedNet(object):
141        """
142            A fully-connected neural network with an arbitrary number of hidden layers,
143            ReLU nonlinearities, and a softmax loss function. This will also implement
144            dropout and batch normalization as options. For a network with L layers,
145            the architecture will be
146
147            {affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax
148
149            where batch normalization and dropout are optional, and the {...} block is
150            repeated L - 1 times.
151
152            Similar to the TwoLayerNet above, learnable parameters are stored in the
153            self.params dictionary and will be learned using the Solver class.
154        """
155
156        def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
157                     dropout=0, use_batchnorm=False, reg=0.0,
158                     weight_scale=1e-2, dtype=np.float32, seed=None):
159        """
160            Initialize a new FullyConnectedNet.
161

```

```

162     Inputs:
163     - hidden_dims: A list of integers giving the size of each hidden layer.
164     - input_dim: An integer giving the size of the input.
165     - num_classes: An integer giving the number of classes to classify.
166     - dropout: Scalar between 0 and 1 giving dropout strength. If dropout=0 then
167       the network should not use dropout at all.
168     - use_batchnorm: Whether or not the network should use batch normalization.
169     - reg: Scalar giving L2 regularization strength.
170     - weight_scale: Scalar giving the standard deviation for random
171       initialization of the weights.
172     - dtype: A numpy datatype object; all computations will be performed using
173       this datatype. float32 is faster but less accurate, so you should use
174       float64 for numeric gradient checking.
175     - seed: If not None, then pass this random seed to the dropout layers. This
176       will make the dropout layers deterministic so we can gradient check the
177       model.
178     """
179     self.use_batchnorm = use_batchnorm
180     self.use_dropout = dropout > 0
181     self.reg = reg
182     self.num_layers = 1 + len(hidden_dims)
183     self.dtype = dtype
184     self.params = {}
185
186     # ===== #
187     # YOUR CODE HERE:
188     #   Initialize all parameters of the network in the self.params dictionary.
189     #   The weights and biases of layer 1 are W1 and b1; and in general the
190     #   weights and biases of layer i are Wi and bi. The
191     #   biases are initialized to zero and the weights are initialized
192     #   so that each parameter has mean 0 and standard deviation weight_scale.
193     # ===== #
194
195     next_layer_input_dim = input_dim
196     for i, hidden_dim in enumerate(hidden_dims, start=1):
197         self.params['W'+str(i)] = weight_scale * np.random.randn(next_layer_input_dim, hidden_dim)
198         self.params['b'+str(i)] = np.zeros(hidden_dim)
199         next_layer_input_dim = hidden_dim
200
201     self.params['W'+str(self.num_layers)] = weight_scale * np.random.randn(next_layer_input_dim, num_classes)
202     self.params['b'+str(self.num_layers)] = np.zeros(num_classes)
203
204     # ===== #
205     # END YOUR CODE HERE
206     # ===== #
207
208     # When using dropout we need to pass a dropout_param dictionary to each
209     # dropout layer so that the layer knows the dropout probability and the mode
210     # (train / test). You can pass the same dropout_param to each dropout layer.
211     self.dropout_param = {}
212     if self.use_dropout:
213         self.dropout_param = {'mode': 'train', 'p': dropout}
214         if seed is not None:
215             self.dropout_param['seed'] = seed
216
217     # With batch normalization we need to keep track of running means and
218     # variances, so we need to pass a special bn_param object to each batch
219     # normalization layer. You should pass self.bn_params[0] to the forward pass
220     # of the first batch normalization layer, self.bn_params[1] to the forward
221     # pass of the second batch normalization layer, etc.
222     self.bn_params = []
223     if self.use_batchnorm:
224         self.bn_params = [{'mode': 'train'} for i in np.arange(self.num_layers - 1)]
225
226     # Cast all parameters to the correct datatype
227     for k, v in self.params.items():
228         self.params[k] = v.astype(dtype)
229
230
231     def loss(self, X, y=None):
232         """
233         Compute loss and gradient for the fully-connected net.
234
235         Input / output: Same as TwoLayerNet above.
236         """
237         X = X.astype(self.dtype)
238         mode = 'test' if y is None else 'train'
239
240         # Set train/test mode for batchnorm params and dropout param since they
241         # behave differently during training and testing.
242         if self.dropout_param is not None:
243             self.dropout_param['mode'] = mode

```

```

244     if self.use_batchnorm:
245         for bn_param in self.bn_params:
246             bn_param[mode] = mode
247
248     scores = None
249
250     # ===== #
251     # YOUR CODE HERE:
252     #   Implement the forward pass of the FC net and store the output
253     #   scores as the variable "scores".
254     # ===== #
255
256     caches = {}
257     layer_out = X
258     for i in range(1, self.num_layers):
259         layer_out, caches[i] = affine_relu_forward(np.array(layer_out),
260                                                 self.params['W'+str(i)],
261                                                 self.params['b'+str(i)])
262
263     scores, caches[self.num_layers] = affine_forward(np.array(layer_out),
264                                                 self.params['W'+str(self.num_layers)],
265                                                 self.params['b'+str(self.num_layers)])
266
267     # ===== #
268     # END YOUR CODE HERE
269     # ===== #
270
271     # If test mode return early
272     if mode == 'test':
273         return scores
274
275     loss, grads = 0.0, {}
276     # ===== #
277     # YOUR CODE HERE:
278     #   Implement the backwards pass of the FC net and store the gradients
279     #   in the grads dict, so that grads[k] is the gradient of self.params[k]
280     #   Be sure your L2 regularization includes a 0.5 factor.
281     # ===== #
282     sum_of_W_norms = 0
283     for i in range(self.num_layers):
284         sum_of_W_norms += (np.linalg.norm(self.params['W'+str(i+1)])) ** 2
285
286     loss, soft_grad = softmax_loss(scores, y)
287     loss += 0.5 * self.reg * sum_of_W_norms
288
289     layer_back_grad, grads['W'+str(self.num_layers)], grads['b'+str(self.num_layers)] = \
290         affine_backward(soft_grad, caches[self.num_layers])
291     grads['W' + str(self.num_layers)] += self.reg * self.params['W' + str(self.num_layers)]
292
293     for i in range(self.num_layers - 1, 0, -1):
294         layer_back_grad, grads['W'+str(i)], grads['b'+str(i)] = affine_relu_backward(layer_back_grad, caches[i])
295         grads['W' + str(i)] += self.reg * self.params['W' + str(i)]
296
297     # ===== #
298     # END YOUR CODE HERE
299     # ===== #
300
301     return loss, grads

```