

## layers.py

```

1 import numpy as np
2 import pdb
3
4 """
5 This code was originally written for CS 231n at Stanford University
6 (cs231n.stanford.edu). It has been modified in various areas for use in the
7 ECE 239AS class at UCLA. This includes the descriptions of what code to
8 implement as well as some slight potential changes in variable names to be
9 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung for
10 permission to use this code. To see the original version, please visit
11 cs231n.stanford.edu.
12 """
13
14
15 def affine_forward(x, w, b):
16     """
17     Computes the forward pass for an affine (fully-connected) layer.
18
19     The input x has shape (N, d_1, ..., d_k) and contains a minibatch of N
20     examples, where each example x[i] has shape (d_1, ..., d_k). We will
21     reshape each input into a vector of dimension D = d_1 * ... * d_k, and
22     then transform it to an output vector of dimension M.
23
24     Inputs:
25     - x: A numpy array containing input data, of shape (N, d_1, ..., d_k)
26     - w: A numpy array of weights, of shape (D, M)
27     - b: A numpy array of biases, of shape (M,)
28
29     Returns a tuple of:
30     - out: output, of shape (N, M)
31     - cache: (x, w, b)
32     """
33
34     # ===== #
35     # YOUR CODE HERE:
36     # Calculate the output of the forward pass. Notice the dimensions
37     # of w are D x M, which is the transpose of what we did in earlier
38     # assignments.
39     # ===== #
40
41     out = x.reshape(x.shape[0], -1).dot(w) + b
42
43     # ===== #
44     # END YOUR CODE HERE
45     # ===== #
46
47     cache = (x, w, b)
48     return out, cache
49
50
51 def affine_backward(dout, cache):
52     """
53     Computes the backward pass for an affine layer.
54
55     Inputs:
56     - dout: Upstream derivative, of shape (N, M)
57     - cache: Tuple of:
58         - x: Input data, of shape (N, d_1, ..., d_k)
59         - w: Weights, of shape (D, M)
60
61     Returns a tuple of:
62     - dx: Gradient with respect to x, of shape (N, d_1, ..., d_k)
63     - dw: Gradient with respect to w, of shape (D, M)
64     - db: Gradient with respect to b, of shape (M,)
65     """
66     x, w, b = cache
67     dx, dw, db = None, None, None
68
69     # ===== #
70     # YOUR CODE HERE:
71     # Calculate the gradients for the backward pass.
72     # ===== #

```

```

73
74     # dout is N x M
75     # dx should be N x d1 x ... x dk; it relates to dout through multiplication with w, which is D x M
76     # dw should be D x M; it relates to dout through multiplication with x, which is N x D after reshaping
77     # db should be M; it is just the sum over dout examples
78
79     db = np.sum(dout, axis=0)
80     dx = dout.dot(w.T).reshape(x.shape)
81     dw = x.reshape(x.shape[0], -1).T.dot(dout)
82     # ===== #
83     # END YOUR CODE HERE
84     # ===== #
85
86     return dx, dw, db
87
88 def relu_forward(x):
89     """
90         Computes the forward pass for a layer of rectified linear units (ReLUs).
91
92         Input:
93         - x: Inputs, of any shape
94
95         Returns a tuple of:
96         - out: Output, of the same shape as x
97         - cache: x
98     """
99     # ===== #
100    # YOUR CODE HERE:
101    # Implement the ReLU forward pass.
102    # ===== #
103
104    out = x * (x > 0)
105    # ===== #
106    # END YOUR CODE HERE
107    # ===== #
108
109    cache = x
110    return out, cache
111
112
113 def relu_backward(dout, cache):
114     """
115         Computes the backward pass for a layer of rectified linear units (ReLUs).
116
117         Input:
118         - dout: Upstream derivatives, of any shape
119         - cache: Input x, of same shape as dout
120
121         Returns:
122         - dx: Gradient with respect to x
123     """
124     x = cache
125
126     # ===== #
127     # YOUR CODE HERE:
128     # Implement the ReLU backward pass
129     # ===== #
130
131     # ReLU directs linearly to those > 0
132     dx = dout * (x > 0)
133
134     # ===== #
135     # END YOUR CODE HERE
136     # ===== #
137
138     return dx
139
140 def svm_loss(x, y):
141     """
142         Computes the loss and gradient using for multiclass SVM classification.
143
144         Inputs:
145         - x: Input data, of shape (N, C) where x[i, j] is the score for the jth class
146             for the ith input.
147         - y: Vector of labels, of shape (N,) where y[i] is the label for x[i] and
148             0 <= y[i] < C

```

```
149
150     Returns a tuple of:
151     - loss: Scalar giving the loss
152     - dx: Gradient of the loss with respect to x
153     """
154
155     N = x.shape[0]
156     correct_class_scores = x[np.arange(N), y]
157     margins = np.maximum(0, x - correct_class_scores[:, np.newaxis] + 1.0)
158     margins[np.arange(N), y] = 0
159     loss = np.sum(margins) / N
160     num_pos = np.sum(margins > 0, axis=1)
161     dx = np.zeros_like(x)
162     dx[margins > 0] = 1
163     dx[np.arange(N), y] -= num_pos
164     dx /= N
165
166
167 def softmax_loss(x, y):
168     """
169     Computes the loss and gradient for softmax classification.
170
171     Inputs:
172     - x: Input data, of shape (N, C) where x[i, j] is the score for the jth class
173       for the ith input.
174     - y: Vector of labels, of shape (N,) where y[i] is the label for x[i] and
175       0 <= y[i] < C
176
177     Returns a tuple of:
178     - loss: Scalar giving the loss
179     - dx: Gradient of the loss with respect to x
180     """
181
182     probs = np.exp(x - np.max(x, axis=1, keepdims=True))
183     probs /= np.sum(probs, axis=1, keepdims=True)
184     N = x.shape[0]
185     loss = -np.sum(np.log(probs[np.arange(N), y])) / N
186     dx = probs.copy()
187     dx[np.arange(N), y] -= 1
188     dx /= N
189
190     return loss, dx
```