

## neural\_net.py

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3 """
4 This code was originally written for CS 231n at Stanford University
5 (cs231n.stanford.edu). It has been modified in various areas for use in the
6 ECE 239AS class at UCLA. This includes the descriptions of what code to
7 implement as well as some slight potential changes in variable names to be
8 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung for
9 permission to use this code. To see the original version, please visit
10 cs231n.stanford.edu.
11 """
12
13
14 class TwoLayerNet(object):
15 """
16 A two-layer fully-connected neural network. The net has an input dimension of
17 N, a hidden layer dimension of H, and performs classification over C classes.
18 We train the network with a softmax loss function and L2 regularization on the
19 weight matrices. The network uses a ReLU nonlinearity after the first fully
20 connected layer.
21
22 In other words, the network has the following architecture:
23
24 input - fully connected layer - ReLU - fully connected layer - softmax
25
26 The outputs of the second fully-connected layer are the scores for each class.
27 """
28
29 def __init__(self, input_size, hidden_size, output_size, std=1e-4):
30 """
31 Initialize the model. Weights are initialized to small random values and
32 biases are initialized to zero. Weights and biases are stored in the
33 variable self.params, which is a dictionary with the following keys:
34
35 W1: First layer weights; has shape (H, D)
36 b1: First layer biases; has shape (H,)
37 W2: Second layer weights; has shape (C, H)
38 b2: Second layer biases; has shape (C,)
39
40 Inputs:
41 - input_size: The dimension D of the input data.
42 - hidden_size: The number of neurons H in the hidden layer.
43 - output_size: The number of classes C.
44 """
45 self.params = {}
46 self.params['W1'] = std * np.random.randn(hidden_size, input_size)
47 self.params['b1'] = np.zeros(hidden_size)
48 self.params['W2'] = std * np.random.randn(output_size, hidden_size)
49 self.params['b2'] = np.zeros(output_size)
50
51
52 def loss(self, X, y=None, reg=0.0):
53 """
54 Compute the loss and gradients for a two layer fully connected neural
55 network.
56
57 Inputs:
58 - X: Input data of shape (N, D). Each X[i] is a training sample.
59 - y: Vector of training labels. y[i] is the label for X[i], and each y[i] is
60 an integer in the range 0 <= y[i] < C. This parameter is optional; if it
61 is not passed then we only return scores, and if it is passed then we
62 instead return the loss and gradients.
63 - reg: Regularization strength.
64
65 Returns:
66 If y is None, return a matrix scores of shape (N, C) where scores[i, c] is
67 the score for class c on input X[i].
68
69 If y is not None, instead return a tuple of:
70 - loss: Loss (data loss and regularization loss) for this batch of training
71 samples.
72 - grads: Dictionary mapping parameter names to gradients of those parameters
73 with respect to the loss function; has the same keys as self.params.
74 """
75 # Unpack variables from the params dictionary
76 W1, b1 = self.params['W1'], self.params['b1']
77 W2, b2 = self.params['W2'], self.params['b2']
78 N, D = X.shape
79
80 # Compute the forward pass
81 scores = None
82

```

```

83      # ===== #
84      # YOUR CODE HERE:
85      # Calculate the output scores of the neural network. The result
86      # should be (N, C). As stated in the description for this class,
87      # there should not be a ReLU layer after the second FC layer.
88      # The output of the second FC layer is the output scores. Do not
89      # use a for loop in your implementation.
90      # ===== #
91
92      def softmax(x):
93          e_x = np.exp(x - np.max(x))
94          return e_x / e_x.sum()
95
96      XdotW1_T = X.dot(W1.T)
97      perceptron_1_out = XdotW1_T + b1
98      layer_1_out = perceptron_1_out * (perceptron_1_out > 0)
99      scores = layer_1_out.dot(W2.T) + b2
100
101     # ===== #
102     # END YOUR CODE HERE
103     # ===== #
104
105
106     # If the targets are not given then jump out, we're done
107     if y is None:
108         return scores
109
110     # Compute the loss
111     loss = None
112
113     # ===== #
114     # YOUR CODE HERE:
115     # Calculate the loss of the neural network. This includes the
116     # softmax loss and the L2 regularization for W1 and W2. Store the
117     # total loss in teh variable loss. Multiply the regularization
118     # loss by 0.5 (in addition to the factor reg).
119     # ===== #
120
121     # scores is num_examples by num_classes
122     W1_norm_sqrd = (1/2) * (np.linalg.norm(W1) ** 2)
123     W2_norm_sqrd = (1/2) * (np.linalg.norm(W2) ** 2)
124     minibatch_size = scores.shape[0]
125     loss = (1 / float(minibatch_size)) * np.sum(
126     np.log(np.sum(np.exp(scores.T), axis=0)) - np.choose(y, scores.T)) + reg * (W1_norm_sqrd + W2_norm_sqrd)
127     # ===== #
128     # END YOUR CODE HERE
129     # ===== #
130
131     grads = {}
132
133     # ===== #
134     # YOUR CODE HERE:
135     # Implement the backward pass. Compute the derivatives of the
136     # weights and the biases. Store the results in the grads
137     # dictionary. e.g., grads['W1'] should store the gradient for
138     # W1, and be of the same size as W1.
139     # ===== #
140
141     """
142     Calculation of softmax gradient with respect to the scores of the last layer. The code is taken from the previous
143     homework, with slight modifications
144     """
145
146     softmax_nominators = np.exp(scores.T - np.amax(scores.T, axis=0))
147     softmax_matrix = softmax_nominators / np.sum(softmax_nominators, axis=0)
148     softmax_matrix[y, np.arange(N)] -= 1
149     softmax_grad = (1 / N) * softmax_matrix
150
151     grads['b2'] = np.sum(softmax_grad, axis=1)
152     grads['W2'] = softmax_grad.dot(layer_1_out) + reg * W2
153     relu_activations = np.array([1] * (perceptron_1_out > 0))
154     grad_bef_relu = W2.T.dot(softmax_grad) * relu_activations.T
155     grads['b1'] = np.sum(grad_bef_relu, axis=1)
156     grads['W1'] = grad_bef_relu.dot(X) + reg * W1
157
158
159     # ===== #
160     # END YOUR CODE HERE
161     # ===== #
162
163     return loss, grads
164
165     def train(self, X, y, X_val, y_val,
166               learning_rate=1e-3, learning_rate_decay=0.95,
167               reg=1e-5, num_iters=100,

```

```

168         batch_size=200, verbose=False):
169     """
170     Train this neural network using stochastic gradient descent.
171
172     Inputs:
173     - X: A numpy array of shape (N, D) giving training data.
174     - y: A numpy array f shape (N,) giving training labels; y[i] = c means that
175         X[i] has label c, where 0 <= c < C.
176     - X_val: A numpy array of shape (N_val, D) giving validation data.
177     - y_val: A numpy array of shape (N_val,) giving validation labels.
178     - learning_rate: Scalar giving learning rate for optimization.
179     - learning_rate_decay: Scalar giving factor used to decay the learning rate
180         after each epoch.
181     - reg: Scalar giving regularization strength.
182     - num_iters: Number of steps to take when optimizing.
183     - batch_size: Number of training examples to use per step.
184     - verbose: boolean; if true print progress during optimization.
185     """
186
187     num_train = X.shape[0]
188     iterations_per_epoch = max(num_train / batch_size, 1)
189
190     # Use SGD to optimize the parameters in self.model
191     loss_history = []
192     train_acc_history = []
193     val_acc_history = []
194
195     for it in np.arange(num_iters):
196         X_batch = None
197         y_batch = None
198
199         # ===== #
200         # YOUR CODE HERE:
201         #   Create a minibatch by sampling batch_size samples randomly.
202         # ===== #
203         idx = np.random.randint(low=0, high=X.shape[0], size=batch_size)
204         X_batch = X[idx]
205         y_batch = y[idx]
206
207         # ===== #
208         # END YOUR CODE HERE
209         # ===== #
210
211         # Compute loss and gradients using the current minibatch
212         loss, grads = self.loss(X_batch, y=y_batch, reg=reg)
213         loss_history.append(loss)
214
215         # ===== #
216         # YOUR CODE HERE:
217         #   Perform a gradient descent step using the minibatch to update
218         #   all parameters (i.e., W1, W2, b1, and b2).
219         # ===== #
220
221         self.params['W1'] -= learning_rate * grads['W1']
222         self.params['b1'] -= learning_rate * grads['b1']
223         self.params['W2'] -= learning_rate * grads['W2']
224         self.params['b2'] -= learning_rate * grads['b2']
225
226         # ===== #
227         # END YOUR CODE HERE
228         # ===== #
229
230         if verbose and it % 100 == 0:
231             print('iteration {} / {}: loss {}'.format(it, num_iters, loss))
232
233         # Every epoch, check train and val accuracy and decay learning rate.
234         if it % iterations_per_epoch == 0:
235             # Check accuracy
236             train_acc = (self.predict(X_batch) == y_batch).mean()
237             val_acc = (self.predict(X_val) == y_val).mean()
238             train_acc_history.append(train_acc)
239             val_acc_history.append(val_acc)
240
241             # Decay learning rate
242             learning_rate *= learning_rate_decay
243
244     return {
245         'loss_history': loss_history,
246         'train_acc_history': train_acc_history,
247         'val_acc_history': val_acc_history,
248     }
249
250     def predict(self, X):
251         """
252         Use the trained weights of this two-layer network to predict labels for
253         data points. For each data point we predict scores for each of the C

```

```
253     classes, and assign each data point to the class with the highest score.  
254  
255     Inputs:  
256     - X: A numpy array of shape (N, D) giving N D-dimensional data points to  
257         classify.  
258  
259     Returns:  
260     - y_pred: A numpy array of shape (N,) giving predicted labels for each of  
261         the elements of X. For all i, y_pred[i] = c means that X[i] is predicted  
262         to have class c, where 0 <= c < C.  
263     """  
264     y_pred = None  
265  
266     # ===== #  
267     # YOUR CODE HERE:  
268     # Predict the class given the input data.  
269     # ===== #  
270     scores = self.loss(X)  
271     y_pred = np.argmax(scores, axis=1)  
272  
273  
274     # ===== #  
275     # END YOUR CODE HERE  
276     # ===== #  
277  
278     return y_pred  
279  
280  
281
```