

fc_net.py

```

1  import numpy as np
2  import pdb
3
4  from .layers import *
5  from .layer_utils import *
6
7  """
8  This code was originally written for CS 231n at Stanford University
9  (cs231n.stanford.edu). It has been modified in various areas for use in the
10 ECE 239AS class at UCLA. This includes the descriptions of what code to
11 implement as well as some slight potential changes in variable names to be
12 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung for
13 permission to use this code. To see the original version, please visit
14 cs231n.stanford.edu.
15 """
16
17 class FullyConnectedNet(object):
18     """
19         A fully-connected neural network with an arbitrary number of hidden layers,
20         ReLU nonlinearities, and a softmax loss function. This will also implement
21         dropout and batch normalization as options. For a network with L layers,
22         the architecture will be
23
24         {affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax
25
26         where batch normalization and dropout are optional, and the {...} block is
27         repeated L - 1 times.
28
29         Similar to the TwoLayerNet above, learnable parameters are stored in the
30         self.params dictionary and will be learned using the Solver class.
31     """
32
33     def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
34                 dropout=0, use_batchnorm=False, reg=0.0,
35                 weight_scale=1e-2, dtype=np.float32, seed=None):
36
37         Initialize a new FullyConnectedNet.
38
39         Inputs:
40         - hidden_dims: A list of integers giving the size of each hidden layer.
41         - input_dim: An integer giving the size of the input.
42         - num_classes: An integer giving the number of classes to classify.
43         - dropout: Scalar between 0 and 1 giving dropout strength. If dropout=1 then
44             the network should not use dropout at all.
45         - use_batchnorm: Whether or not the network should use batch normalization.
46         - reg: Scalar giving L2 regularization strength.
47         - weight_scale: Scalar giving the standard deviation for random
48             initialization of the weights.
49         - dtype: A numpy datatype object; all computations will be performed using
50             this datatype. float32 is faster but less accurate, so you should use
51             float64 for numeric gradient checking.
52         - seed: If not None, then pass this random seed to the dropout layers. This
53             will make the dropout layers deterministic so we can gradient check the
54             model.
55
56         self.use_batchnorm = use_batchnorm
57         self.use_dropout = dropout < 1
58         self.reg = reg
59         self.num_layers = 1 + len(hidden_dims)
60         self.dtype = dtype
61         self.params = {}
62
63         # ===== #
64         # YOUR CODE HERE:
65         #   Initialize all parameters of the network in the self.params dictionary.
66         #   The weights and biases of layer 1 are W1 and b1; and in general the
67         #   weights and biases of layer i are Wi and bi. The
68         #   biases are initialized to zero and the weights are initialized
69         #   so that each parameter has mean 0 and standard deviation weight_scale.
70
71         # BATCHNORM: Initialize the gammas of each layer to 1 and the beta
72         # parameters to zero. The gamma and beta parameters for layer 1 should
73         # be self.params['gammal'] and self.params['betal']. For layer 2, they
74         # should be gamma2 and beta2, etc. Only use batchnorm if self.use_batchnorm
75         # is true and DO NOT do batch normalize the output scores.
76
77         next_layer_input_dim = input_dim
78         for i, hidden_dim in enumerate(hidden_dims, start=1):
79             self.params['W' + str(i)] = weight_scale * np.random.randn(next_layer_input_dim, hidden_dim)
80

```

```

81         self.params['b' + str(i)] = np.zeros(hidden_dim)
82     if self.use_batchnorm:
83         self.params['gamma' + str(i)] = np.ones(hidden_dim)
84         self.params['beta' + str(i)] = np.zeros(hidden_dim)
85     next_layer_input_dim = hidden_dim
86
87     self.params['W' + str(self.num_layers)] = weight_scale * np.random.randn(next_layer_input_dim, num_classes)
88     self.params['b' + str(self.num_layers)] = np.zeros(num_classes)
89
90     # ===== #
91     # END YOUR CODE HERE
92     # ===== #
93
94     # When using dropout we need to pass a dropout_param dictionary to each
95     # dropout layer so that the layer knows the dropout probability and the mode
96     # (train / test). You can pass the same dropout_param to each dropout layer.
97     self.dropout_param = {}
98     if self.use_dropout:
99         self.dropout_param = {'mode': 'train', 'p': dropout}
100    if seed is not None:
101        self.dropout_param['seed'] = seed
102
103    # With batch normalization we need to keep track of running means and
104    # variances, so we need to pass a special bn_param object to each batch
105    # normalization layer. You should pass self.bn_params[0] to the forward pass
106    # of the first batch normalization layer, self.bn_params[1] to the forward
107    # pass of the second batch normalization layer, etc.
108    self.bn_params = []
109    if self.use_batchnorm:
110        self.bn_params = [{ 'mode': 'train'} for i in np.arange(self.num_layers - 1)]
111
112    # Cast all parameters to the correct datatype
113    for k, v in self.params.items():
114        self.params[k] = v.astype(dtype)
115
116
117 def loss(self, X, y=None):
118     """
119     Compute loss and gradient for the fully-connected net.
120
121     Input / output: Same as TwoLayerNet above.
122     """
123     X = X.astype(self.dtype)
124     mode = 'test' if y is None else 'train'
125
126     # Set train/test mode for batchnorm params and dropout param since they
127     # behave differently during training and testing.
128     if self.dropout_param is not None:
129         self.dropout_param['mode'] = mode
130     if self.use_batchnorm:
131         for bn_param in self.bn_params:
132             bn_param[mode] = mode
133
134     scores = None
135
136     # ===== #
137     # YOUR CODE HERE:
138     # Implement the forward pass of the FC net and store the output
139     # scores as the variable "scores".
140     #
141     # BATCHNORM: If self.use_batchnorm is true, insert a bathnorm layer
142     # between the affine_forward and relu_forward layers. You may
143     # also write an affine_batchnorm_relu() function in layer_utils.py.
144     #
145     # DROPOUT: If dropout is non-zero, insert a dropout layer after
146     # every ReLU layer.
147     # ===== #
148
149     caches = {}
150     dropout_caches = {}
151     layer_out = X
152     for i in range(1, self.num_layers):
153         if self.use_batchnorm:
154             layer_out, caches[i] = affine_batchnorm_relu_forward(np.array(layer_out),
155                                                               self.params['W' + str(i)],
156                                                               self.params['b' + str(i)],
157                                                               self.params['gamma' + str(i)],
158                                                               self.params['beta' + str(i)],
159                                                               self.bn_params[i - 1])
160         else: # No batchnorm
161             layer_out, caches[i] = affine_relu_forward(np.array(layer_out),
162                                                       self.params['W' + str(i)],
163                                                       self.params['b' + str(i)])

```

```

164         if self.use_dropout:
165             layer_out, dropout_caches[i] = dropout_forward(layer_out, self.dropout_param)
166
167
168
169     scores, caches[self.num_layers] = affine_forward(np.array(layer_out),
170                                                 self.params['W' + str(self.num_layers)],
171                                                 self.params['b' + str(self.num_layers)])
172
173     # ===== #
174     # END YOUR CODE HERE
175     # ===== #
176
177     # If test mode return early
178     if mode == 'test':
179         return scores
180
181     loss, grads = 0.0, {}
182     # ===== #
183     # YOUR CODE HERE:
184     # Implement the backwards pass of the FC net and store the gradients
185     # in the grads dict, so that grads[k] is the gradient of self.params[k]
186     # Be sure your L2 regularization includes a 0.5 factor.
187     #
188     # BATCHNORM: Incorporate the backward pass of the batchnorm.
189     #
190     # DROPOUT: Incorporate the backward pass of dropout.
191     # ===== #
192
193     sum_of_W_norms = 0
194     for i in range(self.num_layers):
195         sum_of_W_norms += (np.linalg.norm(self.params['W' + str(i + 1)])) ** 2
196
197     loss, soft_grad = softmax_loss(scores, y)
198     loss += 0.5 * self.reg * sum_of_W_norms
199
200     layer_back_grad, grads['W' + str(self.num_layers)], grads['b' + str(self.num_layers)] = \
201         affine_backward(soft_grad, caches[self.num_layers])
202     grads['W' + str(self.num_layers)] += self.reg * self.params['W' + str(self.num_layers)]
203
204     for i in range(self.num_layers - 1, 0, -1):
205         if self.use_dropout:
206             layer_back_grad = dropout_backward(layer_back_grad, dropout_caches[i])
207         if self.use_batchnorm:
208             layer_back_grad, grads['W' + str(i)], grads['b' + str(i)], \
209             grads['gamma' + str(i)], grads['beta' + str(i)] = \
210                 affine_batchnorm_relu_backward(layer_back_grad, caches[i])
211         else:
212             layer_back_grad, grads['W' + str(i)], grads['b' + str(i)] = affine_relu_backward(layer_back_grad,
213                                                 caches[i])
214             grads['W' + str(i)] += self.reg * self.params['W' + str(i)]
215
216     # ===== #
217     # END YOUR CODE HERE
218     # ===== #
219
220     return loss, grads
221

```