

## cnn.py

```

1  import numpy as np
2
3  from nnndl.layers import *
4  from nnndl.conv_layers import *
5  from cs231n.fast_layers import *
6  from nnndl.layer_utils import *
7  from nnndl.conv_layer_utils import *
8
9  import pdb
10
11 """
12 This code was originally written for CS 231n at Stanford University
13 (cs231n.stanford.edu). It has been modified in various areas for use in the
14 ECE 239AS class at UCLA. This includes the descriptions of what code to
15 implement as well as some slight potential changes in variable names to be
16 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung for
17 permission to use this code. To see the original version, please visit
18 cs231n.stanford.edu.
19 """
20
21 class ThreeLayerConvNet(object):
22     """
23     A three-layer convolutional network with the following architecture:
24
25     conv - relu - 2x2 max pool - affine - relu - affine - softmax
26
27     The network operates on minibatches of data that have shape (N, C, H, W)
28     consisting of N images, each with height H and width W and with C input
29     channels.
30     """
31
32     def __init__(self, input_dim=(3, 32, 32), num_filters=32, filter_size=7,
33                  hidden_dim=100, num_classes=10, weight_scale=1e-3, reg=0.0,
34                  dtype=np.float32, use_batchnorm=False):
35         """
36         Initialize a new network.
37
38         Inputs:
39         - input_dim: Tuple (C, H, W) giving size of input data
40         - num_filters: Number of filters to use in the convolutional layer
41         - filter_size: Size of filters to use in the convolutional layer
42         - hidden_dim: Number of units to use in the fully-connected hidden layer
43         - num_classes: Number of scores to produce from the final affine layer.
44         - weight_scale: Scalar giving standard deviation for random initialization
45         of weights.
46         - reg: Scalar giving L2 regularization strength
47         - dtype: numpy datatype to use for computation.
48         """
49         self.use_batchnorm = use_batchnorm
50         self.params = {}
51         self.reg = reg
52         self.dtype = dtype
53
54
55         # ===== #
56         # YOUR CODE HERE:
57         #   Initialize the weights and biases of a three layer CNN. To initialize:
58         #       - the biases should be initialized to zeros.
59         #       - the weights should be initialized to a matrix with entries
60         #           drawn from a Gaussian distribution with zero mean and
61         #           standard deviation given by weight_scale.
62         # ===== #
63
64         self.params['b1'] = np.zeros((num_filters))
65         self.params['b2'] = np.zeros((hidden_dim))
66         self.params['b3'] = np.zeros((num_classes))
67         self.params['W1'] = np.random.normal(0, weight_scale, size=(num_filters, input_dim[0], filter_size, filter_size))
68         self.params['W2'] = np.random.normal(0, weight_scale, size=(int(((input_dim[1] - 2) / 2 + 1) *
69                                         ((input_dim[2] - 2) / 2 + 1) * num_filters),
70                                         hidden_dim))
71         self.params['W3'] = np.random.normal(0, weight_scale, size=(hidden_dim, num_classes))
72         # ===== #
73         # END YOUR CODE HERE
74         # ===== #
75
76         for k, v in self.params.items():
77             self.params[k] = v.astype(dtype)
78
79
80     def loss(self, X, y=None):
81         """

```

```

82     Evaluate loss and gradient for the three-layer convolutional network.
83
84     Input / output: Same API as TwoLayerNet in fc_net.py.
85     """
86     W1, b1 = self.params['W1'], self.params['b1']
87     W2, b2 = self.params['W2'], self.params['b2']
88     W3, b3 = self.params['W3'], self.params['b3']
89
90     # pass conv_param to the forward pass for the convolutional layer
91     filter_size = W1.shape[2]
92     conv_param = {'stride': 1, 'pad': (filter_size - 1) / 2}
93
94     # pass pool_param to the forward pass for the max-pooling layer
95     pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}
96
97     scores = None
98
99     # ===== #
100    # YOUR CODE HERE:
101    #   Implement the forward pass of the three layer CNN. Store the output
102    #   scores as the variable "scores".
103    # ===== #
104    conv_out, conv_cache = conv_relu_pool_forward(X, W1, b1, conv_param, pool_param)
105    affine1_out, affine1_cache = affine_relu_forward(conv_out, W2, b2)
106    scores, affine2_cache = affine_forward(affine1_out, W3, b3)
107    # ===== #
108    # END YOUR CODE HERE
109    # ===== #
110
111    if y is None:
112        return scores
113
114    loss, grads = 0, {}
115    # ===== #
116    # YOUR CODE HERE:
117    #   Implement the backward pass of the three layer CNN. Store the grads
118    #   in the grads dictionary, exactly as before (i.e., the gradient of
119    #   self.params[k] will be grads[k]). Store the loss as "loss", and
120    #   don't forget to add regularization on ALL weight matrices.
121    # ===== #
122    loss, dout = softmax_loss(scores, y)
123    loss += 0.5 * self.reg * ((np.linalg.norm(self.params['W1'] ** 2)) + (np.linalg.norm(self.params['W2'] ** 2)) +
124                                (np.linalg.norm(self.params['W3'] ** 2)))
125    dout, grads['W3'], grads['b3'] = affine_backward(dout, affine2_cache)
126    dout, grads['W2'], grads['b2'] = affine_relu_backward(dout, affine1_cache)
127    dout, grads['W1'], grads['b1'] = conv_relu_pool_backward(dout, conv_cache)
128
129    grads['W3'] += self.reg * self.params['W3']
130    grads['W2'] += self.reg * self.params['W2']
131    grads['W1'] += self.reg * self.params['W1']
132    # ===== #
133    # END YOUR CODE HERE
134    # ===== #
135
136    return loss, grads
137

```