

conv_layers.py

```

1 import numpy as np
2 from nndl.layers import *
3 import pdb
4
5 """
6 This code was originally written for CS 231n at Stanford University
7 (cs231n.stanford.edu). It has been modified in various areas for use in the
8 ECE 239AS class at UCLA. This includes the descriptions of what code to
9 implement as well as some slight potential changes in variable names to be
10 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung for
11 permission to use this code. To see the original version, please visit
12 cs231n.stanford.edu.
13 """
14
15 def conv_forward_naive(x, w, b, conv_param):
16     """
17     A naive implementation of the forward pass for a convolutional layer.
18
19     The input consists of  $N$  data points, each with  $C$  channels, height  $H$  and width
20      $W$ . We convolve each input with  $F$  different filters, where each filter spans
21     all  $C$  channels and has height  $HH$  and width  $WW$ .
22
23     Input:
24     -  $x$ : Input data of shape  $(N, C, H, W)$ 
25     -  $w$ : Filter weights of shape  $(F, C, HH, WW)$ 
26     -  $b$ : Biases, of shape  $(F,)$ 
27     -  $conv\_param$ : A dictionary with the following keys:
28         - 'stride': The number of pixels between adjacent receptive fields in the
29             horizontal and vertical directions.
30         - 'pad': The number of pixels that will be used to zero-pad the input.
31
32     Returns a tuple of:
33     -  $out$ : Output data, of shape  $(N, F, H', W')$  where  $H'$  and  $W'$  are given by
34          $H' = 1 + (H + 2 * pad - HH) / stride$ 
35          $W' = 1 + (W + 2 * pad - WW) / stride$ 
36     -  $cache$ :  $(x, w, b, conv\_param)$ 
37
38     out = None
39     pad = conv_param['pad']
40     stride = conv_param['stride']
41
42     # ===== #
43     # YOUR CODE HERE:
44     # Implement the forward pass of a convolutional neural network.
45     # Store the output as 'out'.
46     # Hint: to pad the array, you can use the function np.pad.
47     # ===== #
48     padded_x = np.pad(x, [(0, 0), (0, 0), (pad, pad), (pad, pad)], mode='constant')
49     out = np.zeros(shape=(x.shape[0], w.shape[0], int(1 + (x.shape[2] + 2 * pad - w.shape[2]) / stride),
50                     int(1 + (x.shape[3] + 2 * pad - w.shape[3]) / stride)))
51     for example in range(x.shape[0]):
52         for f in range(out.shape[1]):
53             for i in range(out.shape[2]):
54                 for j in range(out.shape[3]):
55                     out[example, f, i, j] = b[f] + np.sum(w[f] * padded_x[example, :, i * stride:i * stride + w.shape[2],
56                                                     j * stride:j * stride + w.shape[3]])
57     # ===== #
58     # END YOUR CODE HERE
59     # ===== #
60
61     cache = (x, w, b, conv_param)
62     return out, cache
63
64
65 def conv_backward_naive(dout, cache):
66     """
67     A naive implementation of the backward pass for a convolutional layer.
68
69     Inputs:
70     -  $dout$ : Upstream derivatives.
71     -  $cache$ : A tuple of  $(x, w, b, conv\_param)$  as in conv_forward_naive
72
73     Returns a tuple of:
74     -  $dx$ : Gradient with respect to  $x$ 
75     -  $dw$ : Gradient with respect to  $w$ 
76     -  $db$ : Gradient with respect to  $b$ 
77     """
78     dx, dw, db = None, None, None
79
80     N, F, out_height, out_width = dout.shape
81     x, w, b, conv_param = cache
82
83     stride, pad = [conv_param['stride'], conv_param['pad']]
84     xpad = np.pad(x, ((0, 0), (0, 0), (pad, pad), (pad, pad)), mode='constant')
85     num_filts, _, f_height, f_width = w.shape
86
87     # ===== #

```

```

88 # YOUR CODE HERE:
89 # Implement the backward pass of a convolutional neural network.
90 # Calculate the gradients: dx, dw, and db.
91 # ===== #
92 dw = np.zeros_like(w)
93 dx = np.zeros_like(x)
94 db = np.zeros_like(b)
95 dxdad = np.zeros_like(xpad)
96
97 for f in range(dout.shape[1]): # F
98     for example in range(dout.shape[0]): # N
99         for h_tag in range(dout.shape[2]): # H'
100            for w_tag in range(dout.shape[3]): # W'
101                offset_h = stride * h_tag
102                offset_w = stride * w_tag
103                dw[f] += dout[example, f, h_tag, w_tag] * xpad[example, :, offset_h:offset_h+w.shape[2],
104                                            offset_w:offset_w+w.shape[3]]
105                dxdad[example, :, offset_h:offset_h+w.shape[2],
106                      offset_w:offset_w+w.shape[3]] += dout[example, f, h_tag, w_tag] * w[f]
107 db = np.sum(dout, axis=(0, 2, 3))
108 dx = dxdad[:, :, pad:-pad, pad:-pad] # The padded parameters are not relevant.
109 # ===== #
110 # END YOUR CODE HERE
111 # ===== #
112
113 return dx, dw, db
114
115
116 def max_pool_forward_naive(x, pool_param):
117 """
118 A naive implementation of the forward pass for a max pooling layer.
119
120 Inputs:
121 - x: Input data, of shape (N, C, H, W)
122 - pool_param: dictionary with the following keys:
123     - 'pool_height': The height of each pooling region
124     - 'pool_width': The width of each pooling region
125     - 'stride': The distance between adjacent pooling regions
126
127 Returns a tuple of:
128 - out: Output data
129 - cache: (x, pool_param)
130 """
131 out = None
132
133 # ===== #
134 # YOUR CODE HERE:
135 # Implement the max pooling forward pass.
136 # ===== #
137 pool_height = pool_param['pool_height']
138 pool_width = pool_param['pool_width']
139 stride = pool_param['stride']
140
141 out_height = int((x.shape[2] - pool_height) / stride) + 1
142 out_width = int((x.shape[3] - pool_width) / stride) + 1
143 out = np.zeros(shape=(x.shape[0], x.shape[1], out_height, out_width))
144
145 for example in range(out.shape[0]):
146     for c in range(out.shape[1]):
147         for h in range(out.shape[2]):
148             for w in range(out.shape[3]):
149                 out[example, c, h, w] = \
150                     np.amax(x[example, c, h * stride:h * stride + pool_height, w * stride:w * stride + pool_width])
151 # ===== #
152 # END YOUR CODE HERE
153 # ===== #
154 cache = (x, pool_param)
155 return out, cache
156
157 def max_pool_backward_naive(dout, cache):
158 """
159 A naive implementation of the backward pass for a max pooling layer.
160
161 Inputs:
162 - dout: Upstream derivatives
163 - cache: A tuple of (x, pool_param) as in the forward pass.
164
165 Returns:
166 - dx: Gradient with respect to x
167 """
168 dx = None
169 x, pool_param = cache
170 pool_height, pool_width, stride = pool_param['pool_height'], pool_param['pool_width'], pool_param['stride']
171
172 # ===== #
173 # YOUR CODE HERE:
174 # Implement the max pooling backward pass.
175 # ===== #
176 dx = np.zeros_like(x)
177 for f in range(dout.shape[1]): # F
178     for example in range(dout.shape[0]): # N

```

```

179     for h_tag in range(dout.shape[2]): # H'
180         for w_tag in range(dout.shape[3]): # W'
181             pool_window = x[example, f, h_tag * stride:h_tag * stride + pool_height, w_tag * stride:w_tag * stride + pool_width]
182             h_max_index, w_max_index = np.unravel_index(np.argmax(pool_window, axis=None), pool_window.shape)
183             dx[example, f, h_max_index + h_tag*stride, w_max_index + w_tag*stride] = dout[example, f, h_tag, w_tag]
184
185     # ===== #
186     # END YOUR CODE HERE
187     # ===== #
188
189     return dx
190
191 def spatial_batchnorm_forward(x, gamma, beta, bn_param):
192     """
193     Computes the forward pass for spatial batch normalization.
194
195     Inputs:
196     - x: Input data of shape (N, C, H, W)
197     - gamma: Scale parameter, of shape (C,)
198     - beta: Shift parameter, of shape (C,)
199     - bn_param: Dictionary with the following keys:
200         - mode: 'train' or 'test'; required
201         - eps: Constant for numeric stability
202         - momentum: Constant for running mean / variance. momentum=0 means that
203             old information is discarded completely at every time step, while
204             momentum=1 means that new information is never incorporated. The
205             default of momentum=0.9 should work well in most situations.
206         - running_mean: Array of shape (D,) giving running mean of features
207         - running_var Array of shape (D,) giving running variance of features
208
209     Returns a tuple of:
210     - out: Output data, of shape (N, C, H, W)
211     - cache: Values needed for the backward pass
212     """
213     out, cache = None, None
214
215     # ===== #
216     # YOUR CODE HERE:
217     # Implement the spatial batchnorm forward pass.
218     #
219     # You may find it useful to use the batchnorm forward pass you
220     # implemented in HW #4.
221     # ===== #
222     x_hat = x.swapaxes(0, 1).reshape(x.shape[1], -1).T
223     out, cache = np.array(batchnorm_forward(x_hat, gamma, beta, bn_param))
224     out = out.reshape(x.shape[0], x.shape[2], x.shape[3], -1).swapaxes(1, 3).swapaxes(2, 3)
225
226     # ===== #
227     # END YOUR CODE HERE
228     # ===== #
229
230     return out, cache
231
232
233 def spatial_batchnorm_backward(dout, cache):
234     """
235     Computes the backward pass for spatial batch normalization.
236
237     Inputs:
238     - dout: Upstream derivatives, of shape (N, C, H, W)
239     - cache: Values from the forward pass
240
241     Returns a tuple of:
242     - dx: Gradient with respect to inputs, of shape (N, C, H, W)
243     - dgamma: Gradient with respect to scale parameter, of shape (C,)
244     - dbeta: Gradient with respect to shift parameter, of shape (C,)
245     """
246     dx, dgamma, dbeta = None, None, None
247
248     # ===== #
249     # YOUR CODE HERE:
250     # Implement the spatial batchnorm backward pass.
251     #
252     # You may find it useful to use the batchnorm forward pass you
253     # implemented in HW #4.
254     # ===== #
255
256     dout_hat = dout.swapaxes(0, 1).reshape(dout.shape[1], -1).T
257     dx, dgamma, dbeta = np.array(batchnorm_backward(dout_hat, cache))
258     dx = dx.reshape(dout.shape[0], dout.shape[2], dout.shape[3], -1).swapaxes(1, 3).swapaxes(2, 3)
259
260     # ===== #
261     # END YOUR CODE HERE
262     # ===== #
263
264     return dx, dgamma, dbeta

```