

layers.py

```

1  import numpy as np
2  import pdb
3
4  """
5  This code was originally written for CS 231n at Stanford University
6  (cs231n.stanford.edu). It has been modified in various areas for use in the
7  ECE 239AS class at UCLA. This includes the descriptions of what code to
8  implement as well as some slight potential changes in variable names to be
9  consistent with class nomenclature. We thank Justin Johnson & Serena Yeung for
10 permission to use this code. To see the original version, please visit
11 cs231n.stanford.edu.
12 """
13
14 def affine_forward(x, w, b):
15     """
16     Computes the forward pass for an affine (fully-connected) layer.
17
18     The input x has shape (N, d_1, ..., d_k) and contains a minibatch of N
19     examples, where each example x[i] has shape (d_1, ..., d_k). We will
20     reshape each input into a vector of dimension D = d_1 * ... * d_k, and
21     then transform it to an output vector of dimension M.
22
23     Inputs:
24     - x: A numpy array containing input data, of shape (N, d_1, ..., d_k)
25     - w: A numpy array of weights, of shape (D, M)
26     - b: A numpy array of biases, of shape (M,)
27
28     Returns a tuple of:
29     - out: output, of shape (N, M)
30     - cache: (x, w, b)
31     """
32
33     # ===== #
34     # YOUR CODE HERE:
35     # Calculate the output of the forward pass. Notice the dimensions
36     # of w are D x M, which is the transpose of what we did in earlier
37     # assignments.
38     # ===== #
39
40     out = x.reshape(x.shape[0], -1).dot(w) + b
41     # ===== #
42     # END YOUR CODE HERE
43     # ===== #
44
45     cache = (x, w, b)
46     return out, cache
47
48
49 def affine_backward(dout, cache):
50     """
51     Computes the backward pass for an affine layer.
52
53     Inputs:
54     - dout: Upstream derivative, of shape (N, M)
55     - cache: Tuple of:
56         - x: Input data, of shape (N, d_1, ..., d_k)
57         - w: Weights, of shape (D, M)
58
59     Returns a tuple of:
60     - dx: Gradient with respect to x, of shape (N, d_1, ..., d_k)
61     - dw: Gradient with respect to w, of shape (D, M)
62     - db: Gradient with respect to b, of shape (M,)
63     """
64     x, w, b = cache
65     dx, dw, db = None, None, None
66
67     # ===== #
68     # YOUR CODE HERE:
69     # Calculate the gradients for the backward pass.
70     # ===== #
71
72     db = np.sum(dout, axis=0)
73     dx = np.array(dout).dot(w.T).reshape(x.shape)
74     dw = x.reshape(x.shape[0], -1).T.dot(dout)

```

```

75      # ===== #
76      # END YOUR CODE HERE
77      # ===== #
78
79      return dx, dw, db
80
81  def relu_forward(x):
82      """
83          Computes the forward pass for a layer of rectified linear units (ReLUs).
84
85          Input:
86          - x: Inputs, of any shape
87
88          Returns a tuple of:
89          - out: Output, of the same shape as x
90          - cache: x
91      """
92      # ===== #
93      # YOUR CODE HERE:
94      # Implement the ReLU forward pass.
95      # ===== #
96      out = x * (x > 0)
97      # ===== #
98      # END YOUR CODE HERE
99      # ===== #
100
101     cache = x
102     return out, cache
103
104
105  def relu_backward(dout, cache):
106      """
107          Computes the backward pass for a layer of rectified linear units (ReLUs).
108
109          Input:
110          - dout: Upstream derivatives, of any shape
111          - cache: Input x, of same shape as dout
112
113          Returns:
114          - dx: Gradient with respect to x
115      """
116     x = cache
117
118     # ===== #
119     # YOUR CODE HERE:
120     # Implement the ReLU backward pass
121     # ===== #
122     dx = dout * (x > 0)
123
124     # ===== #
125     # END YOUR CODE HERE
126     # ===== #
127
128     return dx
129
130  def batchnorm_forward(x, gamma, beta, bn_param):
131      """
132          Forward pass for batch normalization.
133
134          During training the sample mean and (uncorrected) sample variance are
135          computed from minibatch statistics and used to normalize the incoming data.
136          During training we also keep an exponentially decaying running mean of the mean
137          and variance of each feature, and these averages are used to normalize data
138          at test-time.
139
140          At each timestep we update the running averages for mean and variance using
141          an exponential decay based on the momentum parameter:
142
143          running_mean = momentum * running_mean + (1 - momentum) * sample_mean
144          running_var = momentum * running_var + (1 - momentum) * sample_var
145
146          Note that the batch normalization paper suggests a different test-time
147          behavior: they compute sample mean and variance for each feature using a
148          large number of training images rather than using a running average. For
149          this implementation we have chosen to use running averages instead since
150          they do not require an additional estimation step; the torch7 implementation
151          of batch normalization also uses running averages.

```

```

152
153     Input:
154     - x: Data of shape (N, D)
155     - gamma: Scale parameter of shape (D,)
156     - beta: Shift parameter of shape (D,)
157     - bn_param: Dictionary with the following keys:
158         - mode: 'train' or 'test'; required
159         - eps: Constant for numeric stability
160         - momentum: Constant for running mean / variance.
161         - running_mean: Array of shape (D,) giving running mean of features
162         - running_var Array of shape (D,) giving running variance of features
163
164     Returns a tuple of:
165     - out: of shape (N, D)
166     - cache: A tuple of values needed in the backward pass
167     """
168     mode = bn_param['mode']
169     eps = bn_param.get('eps', 1e-5)
170     momentum = bn_param.get('momentum', 0.9)
171
172     N, D = x.shape
173     running_mean = bn_param.get('running_mean', np.zeros(D, dtype=x.dtype))
174     running_var = bn_param.get('running_var', np.zeros(D, dtype=x.dtype))
175
176     out, cache = None, None
177     if mode == 'train':
178
179         # ===== #
180         # YOUR CODE HERE:
181         #   A few steps here:
182         #       (1) Calculate the running mean and variance of the minibatch.
183         #       (2) Normalize the activations with the running mean and variance.
184         #       (3) Scale and shift the normalized activations. Store this
185         #           as the variable 'out'
186         #       (4) Store any variables you may need for the backward pass in
187         #           the 'cache' variable.
188         # ===== #
189         running_mean = np.mean(x, axis=0)
190         running_var = np.var(x, axis=0)
191         x_hat = (x - running_mean) / np.sqrt(eps + running_var)
192         out = gamma * x_hat + beta
193         cache = (x_hat, x, running_mean, running_var, eps, gamma)
194
195         # ===== #
196         # END YOUR CODE HERE
197         # ===== #
198
199     elif mode == 'test':
200
201         # ===== #
202         # YOUR CODE HERE:
203         #   Calculate the testing time normalized activation. Normalize using
204         #   the running mean and variance, and then scale and shift appropriately.
205         #   Store the output as 'out'.
206         # ===== #
207         x_hat = (x - running_mean) / np.sqrt(eps + running_var)
208         out = gamma * x_hat + beta
209
210         # ===== #
211         # END YOUR CODE HERE
212         # ===== #
213
214     else:
215         raise ValueError('Invalid forward batchnorm mode "%s"' % mode)
216
217     # Store the updated running means back into bn_param
218     bn_param['running_mean'] = running_mean
219     bn_param['running_var'] = running_var
220
221     return out, cache
222
223 def batchnorm_backward(dout, cache):
224     """
225     Backward pass for batch normalization.
226
227     For this implementation, you should write out a computation graph for
228     batch normalization on paper and propagate gradients backward through

```

```

229     intermediate nodes.
230
231     Inputs:
232     - dout: Upstream derivatives, of shape (N, D)
233     - cache: Variable of intermediates from batchnorm_forward.
234
235     Returns a tuple of:
236     - dx: Gradient with respect to inputs x, of shape (N, D)
237     - dgamma: Gradient with respect to scale parameter gamma, of shape (D,)
238     - dbeta: Gradient with respect to shift parameter beta, of shape (D,)
239     """
240     dx, dgamma, dbeta = None, None, None
241
242     # ===== #
243     # YOUR CODE HERE:
244     #   Implement the batchnorm backward pass, calculating dx, dgamma, and dbeta.
245     # ===== #
246     x_hat, x, mean, var, eps, gamma = cache
247     N, D = x.shape
248     dbeta = np.sum(dout, axis=0)
249     dgamma = np.sum(dout * x_hat, axis=0)
250     dl_dxhat = dout * gamma
251     dl_dvar = (-1 / 2) * np.sum((1 / ((var + eps) ** (3 / 2))) * (x - mean) * dl_dxhat, axis=0)
252     dl_dmean = (-1 / (np.sqrt(var + eps))) * np.sum(dl_dxhat, axis=0)
253     dx = np.array((1 / np.sqrt(var + eps)) * dl_dxhat + (2 * (x - mean) / N) * dl_dvar + (1 / N) * dl_dmean)
254
255     # ===== #
256     # END YOUR CODE HERE
257     # ===== #
258
259     return dx, dgamma, dbeta
260
261 def dropout_forward(x, dropout_param):
262     """
263     Performs the forward pass for (inverted) dropout.
264
265     Inputs:
266     - x: Input data, of any shape
267     - dropout_param: A dictionary with the following keys:
268         - p: Dropout parameter. We drop each neuron output with probability p.
269         - mode: 'test' or 'train'. If the mode is train, then perform dropout;
270             if the mode is test, then just return the input.
271         - seed: Seed for the random number generator. Passing seed makes this
272             function deterministic, which is needed for gradient checking but not in
273             real networks.
274
275     Outputs:
276     - out: Array of the same shape as x.
277     - cache: A tuple (dropout_param, mask). In training mode, mask is the dropout
278         mask that was used to multiply the input; in test mode, mask is None.
279     """
280     p, mode = dropout_param['p'], dropout_param['mode']
281     if 'seed' in dropout_param:
282         np.random.seed(dropout_param['seed'])
283
284     mask = None
285     out = None
286
287     if mode == 'train':
288         # ===== #
289         # YOUR CODE HERE:
290         #   Implement the inverted dropout forward pass during training time.
291         #   Store the masked and scaled activations in out, and store the
292         #   dropout mask as the variable mask.
293         # ===== #
294     mask = np.random.rand(*x.shape) < p
295     out = x * mask / p
296
297     # ===== #
298     # END YOUR CODE HERE
299     # ===== #
300
301     elif mode == 'test':
302
303         # ===== #
304         # YOUR CODE HERE:
305         #   Implement the inverted dropout forward pass during test time.

```

```

306     # ===== #
307     out = x
308
309     # ===== #
310     # END YOUR CODE HERE
311     # ===== #
312
313     cache = (dropout_param, mask)
314     out = out.astype(x.dtype, copy=False)
315
316     return out, cache
317
318 def dropout_backward(dout, cache):
319     """
320     Perform the backward pass for (inverted) dropout.
321
322     Inputs:
323     - dout: Upstream derivatives, of any shape
324     - cache: (dropout_param, mask) from dropout_forward.
325     """
326     dropout_param, mask = cache
327     mode = dropout_param['mode']
328
329     dx = None
330     if mode == 'train':
331         # ===== #
332         # YOUR CODE HERE:
333         # Implement the inverted dropout backward pass during training time.
334         # ===== #
335         dx = dout * mask / dropout_param['p']
336
337         # ===== #
338         # END YOUR CODE HERE
339         # ===== #
340     elif mode == 'test':
341         # ===== #
342         # YOUR CODE HERE:
343         # Implement the inverted dropout backward pass during test time.
344         # ===== #
345         dx = dout
346
347         # ===== #
348         # END YOUR CODE HERE
349         # ===== #
350     return dx
351
352 def svm_loss(x, y):
353     """
354     Computes the loss and gradient using for multiclass SVM classification.
355
356     Inputs:
357     - x: Input data, of shape (N, C) where x[i, j] is the score for the jth class
358       for the ith input.
359     - y: Vector of labels, of shape (N,) where y[i] is the label for x[i] and
360       0 <= y[i] < C
361
362     Returns a tuple of:
363     - loss: Scalar giving the loss
364     - dx: Gradient of the loss with respect to x
365     """
366     N = x.shape[0]
367     correct_class_scores = x[np.arange(N), y]
368     margins = np.maximum(0, x - correct_class_scores[:, np.newaxis] + 1.0)
369     margins[np.arange(N), y] = 0
370     loss = np.sum(margins) / N
371     num_pos = np.sum(margins > 0, axis=1)
372     dx = np.zeros_like(x)
373     dx[margins > 0] = 1
374     dx[np.arange(N), y] -= num_pos
375     dx /= N
376     return loss, dx
377
378 def softmax_loss(x, y):
379     """
380     Computes the loss and gradient for softmax classification.
381
382     Inputs:

```

```
383     - x: Input data, of shape (N, C) where x[i, j] is the score for the jth class  
384     for the ith input.  
385     - y: Vector of labels, of shape (N,) where y[i] is the label for x[i] and  
386     0 <= y[i] < C  
387  
388     Returns a tuple of:  
389     - loss: Scalar giving the loss  
390     - dx: Gradient of the loss with respect to x  
391     """  
392  
393     probs = np.exp(x - np.max(x, axis=1, keepdims=True))  
394     probs /= np.sum(probs, axis=1, keepdims=True)  
395     N = x.shape[0]  
396     loss = -np.sum(np.log(probs[np.arange(N), y])) / N  
397     dx = probs.copy()  
398     dx[np.arange(N), y] -= 1  
399     dx /= N  
400     return loss, dx  
401
```